

TalentLoop

Hiring that can explain itself.

An agentic hiring assistant that scores candidates against a transparent rubric, cites its evidence, and never contacts a person without a human saying yes.

TEAM 000

Pranav Somnath Undirkalle · Vineet Ramayya Polampalli

Raj Gautam · Priyanshu Kumar

LIVE NOW

talentloop-rouge.vercel.app

talentloop-api.onrender.com/docs

github.com/FasterThanAi/talentloop

youtu.be/E8CvX2cyF_4

Two people lose in every hiring round

And it is the same broken step that fails both

THE RECRUITER

One posting draws hundreds of applicants. Screening stops long before the list does, so the decision is made by whoever was read first — not by who was strongest.

Keyword filters reward people who write the right words, not people who did the work.

THE CANDIDATE

Applies, waits, and hears nothing. No reason, no signal, nothing to act on — and no way to find out whether the gap was skills, evidence, or timing.

The next application is filed just as blindly as the last one.

The screening step is where both failures happen. That is the step we rebuilt.

Most “AI recruiting” makes it worse

Faster rejection is not better hiring

1

A number with no reasoning

A model returns 87/100. Nobody can say why, reproduce it, or defend it to a candidate who asks.

2

Confident invention

Models fill gaps. An unverified claim about someone's experience is not a screening signal — it is a liability.

3

Automation pointed at people


Bulk outreach and auto-rejection scale the least human part of hiring and remove the one place judgment mattered.

Our position

Better-evidenced automation, not more automation.

How the loop works

Five stages. A human owns the two that touch a person.




Invariant 1 — the model never produces the score

It rates dimensions with citations. Deterministic Python does the arithmetic.




Change a weight in a rubric file and every score recomputes.
Disagreeing with the system does not require re-prompting a black box.



Every dimension carries a justification and verified source URLs.

Invariant 2 — nothing reaches a person unapproved

Approval is a separate, audited action. There is no bypass in the codebase.



- **Send is disabled until approved**
Draft and approved are different states, enforced server-side.
- **Do-not-contact is re-checked**
Consent is verified again immediately before dispatch.
- **Every action is audited**
Actor, action, entity and payload, in the same transaction.
- **CI blocks a bypass**
A send path that skips the guard fails the build.

Rejected candidates are told why.

Silence is a product decision. We made the other one.

The part no competitor ships

A rejected candidate leaves with something useful

Every scored candidate gets a report generated from the same breakdown the recruiter saw: what was strong, what evidence was missing, and what to do about it.



Grounded, not generic

Each claim is checked against the score breakdown it came from.




Released deliberately

The recruiter decides when it goes out. Nothing auto-rejects.



Portable

The candidate keeps a verifiable record of the evaluation.



Candidate portal — feedback report

Fairness is a build gate, not a claim on a slide

Every push runs these. A prompt or rubric change that moves a matched-pair probe past tolerance fails the build and cannot be merged.

Bias probes

Matched pairs differing only by identity must score the same.

Sanitised input

Name, photo, age, gender and school never enter the scoring payload.

Citation check

URLs absent from stored evidence are stripped before display.

Feedback fidelity

No report may claim anything the breakdown does not support.

Invariant 1 test

The scoring prompt may never ask for an overall score.

Invariant 2 test

No send path may bypass the approval guard.

Six checks. All blocking. Green on every commit in the public repository.

Built to survive a live demo

The unglamorous engineering that decides whether a product works

Provider failover

Gemini is primary; a quota error fails over to Groq instantly rather than burning retries on a provider that already said no.

Nothing blocks the loop

Every blocking SDK, DNS and HTTP call runs off the event loop, so one slow upstream cannot freeze the API.

Honest degradation

With no key configured the system serves canned data and says so — in the logs, in /health, and as a banner. Demoing on mocks by accident is impossible.


Append-only audit

Action and audit row commit or roll back together, so the trail cannot drift from what actually happened.


React · Vite · Tailwind FastAPI · SQLAlchemy · Alembic Postgres + pgvector Render · Vercel · Supabase

What is live today

Not a prototype behind a login you cannot reach



Requisitions and roles, with pipeline counts



Ranked pipeline — every score opens its evidence

LIVE APP

talentloop-rouge.vercel.app

API DOCS

talentloop-api.onrender.com/docs

REPOSITORY

github.com/FasterThanAi/talentloop

WALKTHROUGH

youtu.be/E8CvX2cyF_4

Demo logins are in the README. Recruiter and candidate see different applications.

What it does not do

Stated plainly, because a demo that hides its edges is not evidence

- **It does not find people's email addresses.**

URL-sourced candidates get a deliberately undeliverable placeholder. A real address arrives only from the recruiter's own CSV or the RizeOS pool.

- **It does not scrape restricted platforms.**

Enrichment reads the public pages a recruiter supplies. Nothing whose terms forbid it.

- **It does not score what it cannot evidence.**

A candidate with no public work gets a low-confidence score, labelled as such rather than dressed up as certainty.

- **It does not decide.**

The rubric ranks and explains. Every action that touches a person still needs a human.

Every one of these is a boundary we chose, not a feature we ran out of time for.

Where it goes next

Ordered by what makes the evaluation more trustworthy, not flashier

1

NOW

Live MVP: parse, source, evidence, explainable scoring, approval-gated outreach, reply triage, candidate feedback, audit trail.

2

NEXT 30 DAYS

Recruiter-tunable rubric weights per requisition. Structured interview questions generated from the specific gaps in a candidate's evidence.

3

60–90 DAYS

Calibration against real hiring outcomes, verifiable feedback credentials for candidates, and an ATS integration so it fits an existing workflow.

The thing worth building on

An evaluation that can be audited, corrected and explained is the only kind that survives contact with a real hiring decision — or a regulator.

A human on every decision that touches a person.

And a candidate who leaves with something they can use.

talentloop-rouge.vercel.app

github.com/FasterThanAi/talentloop · youtu.be/E8CvX2cyF_4

RizeOS Hackathon 2026 · Round 3 — Final Product Showcase

TEAM 000

Pranav Somnath Undirkalle

Backend, data model, docs

Vineet Ramayya Polampalli

Frontend, design system, portal

Raj Gautam

AI pipeline, prompts, evals

Priyanshu Kumar

Architecture, deployment, CI